

Exercise 1: Linear classifier (numpy)

In this exercise, we will learn to implement a Linear classifier using numpy modules. The goal here is to show the basic concepts of classification and to learn how to implement forward/backward propagation of operations used in Linear classifier.

You will be asked to complete several functions/classes used below in `HW_YourAnswer.py`.

- `softmax`
- `cross_entropy_loss`
- `linear_predict`
- `linear_cost_func`
- `batch_gradient_descent_func`
- `stochastic_gradient_descent_func`

First, let's check if you are properly using GPU

- Output should be 'True'
- If not, please follow the instructions in ETL.

```
In [1]: import torch
use_cuda = torch.cuda.is_available()
print('GPU available?:', use_cuda)
```

GPU available?: True

Next, let's mount your drive directory to current notebook and change the system directory to your working directory

```
In [3]: from google.colab import drive
drive.mount('/content/drive')

import sys
sys.path.append('/content/drive/My Drive/Colab Notebooks/Intro_dl/hw1')

import os
os.chdir('/content/drive/My Drive/Colab Notebooks/Intro_dl/hw1')
```

Mounted at /content/drive

```
In [4]: %load_ext autoreload
%autoreload 2

from IPython.display import Image
from utils import *
from HW_YourAnswer import *

%matplotlib inline
```

The autoreload extension is already loaded. To reload it, use:
`%reload_ext autoreload`

Random Seed

```
In [5]: seed = 0
np.random.seed(seed)
```

Data preprocessing

We will use the CIFAR10 dataset.

CIFAR10 dataset is composed of 50000 train data and 10000 test data but we will only use 10000, 1000 images each.

```
In [6]: # 10000 images for training data and 1000 images for test data
# 'get_cifar10_data()' returns train/test data with its mean subtracted for training
X_tr, Y_tr, X_te, Y_te, mean_img = get_CIFAR10_data()
print ('Train data shape : %s, Train labels shape : %s' % (X_tr.shape, Y_tr.shape))
print ('Test data shape : %s, Test labels shape : %s' % (X_te.shape, Y_te.shape))

nb_classes = 10

# one-hot encode labels from integers(0, 1, .., 9) to 10-dim one-hot vector
train_targets = Y_tr.reshape(-1)
Y_tr_onehot = np.eye(nb_classes)[train_targets]

test_targets = Y_te.reshape(-1)
Y_te_onehot = np.eye(nb_classes)[test_targets]
```

```
Train data shape : (10000, 3072), Train labels shape : (10000,)
Test data shape : (1000, 3072), Test labels shape : (1000,)
```

Visualize training images

```
In [6]: # list of classes for cifar10
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

# random data idx for visualization
images_index = np.int32(np.round(np.random.rand(18,)*10000,0))

fig, axes = plt.subplots(3, 6, figsize=(18, 6),
                        subplot_kw={'xticks': [], 'yticks': []})

fig.subplots_adjust(hspace=0.3, wspace=0.05)

for ax, idx in zip(axes.flat, images_index):
    # mean is added again to restore original image
    img = (X_tr[idx,:3072].reshape(32, 32, 3) + mean_img.reshape(32, 32, 3))
    ax.imshow(img)
    ax.set_title(class_names[Y_tr[idx]])
```



1. Softmax function

To do:

- Implement `softmax` in `HW_YourAnswer.py` file

$$p(\mathbf{Y} = k | \mathbf{X} = x) = \frac{\exp(\mathbf{s}_k)}{\sum_{j=1}^K \exp(s_j)}$$

(p.21 Leture 4)

- The output should be `[2.06115362e-09 1.80485138e-35 9.99999998e-01]`
- Sum of the softmax output should be `0.9999999999999999`
- It should not be NaN

```
In [7]: temp_x = np.array([[2060,2000,2080]])
softmax_result1 = softmax(temp_x)
print('Softmax result :\n',softmax_result1)
print ('\nSum of the softmax :',np.sum(softmax_result1))
```

```
Softmax result :
[[2.06115362e-09 1.80485138e-35 9.99999998e-01]]
```

```
Sum of the softmax : 0.9999999999999999
```

- The output should be
`[2.06115362e-09 1.80485138e-35 9.99999998e-01]`
`[2.06106005e-09 4.53978686e-05 9.99954600e-01]`
- Sum of the softmax output should be `[1. 1.]`

```
In [8]: temp_x = np.array([[2060,2000,2080],[1010,1020,1030]])
softmax_result2 = softmax(temp_x)
print('Softmax result :\n',softmax_result2)
print ('\nSum of the softmax :',np.sum(softmax_result2,axis=1))
```

```
Softmax result :
```

```
[[2.06115362e-09 1.80485138e-35 9.99999998e-01]  
[2.06106005e-09 4.53978686e-05 9.99954600e-01]]
```

Sum of the softmax : [1. 1.]

2. Cross-Entropy Loss

Here, we will consider cross-entropy loss between the true target value(e.g., 1) and predicted value after softmax. Note that Negative log-likelihood for the true class and cross entropy between true target and prediction has a similar formula (Lecture 4 p.23).

To do:

- Implement `cross_entropy_loss` in `HW_YourAnswer.py` file

$$L(W) = \frac{1}{N} \sum_{i=1}^N L_i(s_i, y_i)$$

where

$$L_i(s_i, y_i) = \frac{\exp(s_{y_i})}{\sum_{j=1}^K \exp(s_j)}$$

N : number of data

s_i : softmax score for i^{th} class

NOTE: cross entropy loss is **averaged** w.r.t the number of data

- The output should be 20.72326583694641
- It should not be NaN

```
In [9]: temp_score0 = np.array([[0.0, 0.0, 0.0]])  
temp_target0 = np.array([[0,1,0]])  
loss0 = cross_entropy_loss(temp_score0, temp_target0)  
print('Total Loss for temp_0 =', loss0)
```

Total Loss for temp_0 = 20.72326583694641

- The output should be 1.2039728009926025
- It should not be NaN

```
In [10]: temp_score1 = np.array([[0.1, 0.3, 0.6]])  
temp_target1 = np.array([[0,1,0]])  
loss1 = cross_entropy_loss(temp_score1, temp_target1)  
print('Total Loss for temp_1 =', loss1)
```

Total Loss for temp_1 = 1.2039728009926025

- The output should be 0.7418746816378242
- It should not be NaN

```
temp_score2 = np.array([[0.1, 0.3, 0.6], [0.2,0.4,0.4], [0.9,0.05,0.05]])
```

```
In [11]: temp_target2 = np.array([[0,1,0],[0,0,1],[1,0,0]])
loss2 = cross_entropy_loss(temp_score2, temp_target2)
print('Total Loss for temp_2 =', loss2)
```

Total Loss for temp_2 = 0.7418746816378242

3. Linear classifier

With several functions we've defined above, we will implement Linear classifier learned in Lecture 5 p.31-35.

One thing to note is that the **column** of weight matrix indicates weight for each class, which is different from notations used in the Lecture slides.


W_i : i^{th} column of the matrix = classifier for class i

So, the score(before softmax) will be computed as

$$XW + b$$

```
In [12]: Image('fig/Affine_Layer_pic.png')
```

Out[12]:



$$X : (N, D)$$

$$W : (D, C)$$

where N, D, C represents number of data, data dimension, and number of classes, respectively

To do:

- Implement `linear_predict`, `linear_cost_func`, `batch_gradient_descent_func` in `HW_YourAnswer.py` file.

```
In [13]: # Use only 20 data for example
x_batch = X_tr[:20]
y_batch = Y_tr_onehot[:20]

print ('Train data shape : %s, Train labels shape : %s' % (x_batch.shape, y_batch.shape))

# Initialize weights & bias
W = np.zeros((3072, 10))
```

```
b = np.zeros((10))
```

```
# Parameters for the gradient descent
```

```
iterations = 100
```

```
alpha = 0.1
```

Train data shape : (20, 3072), Train labels shape : (20, 10)

- cost should be 2.302585082994045

```
In [14]: initial_cost = linear_cost_func(x_batch, y_batch, w, b)
print('loss:', initial_cost)
```

loss: 2.302585082994045

- final loss should be 0.018715975170939064

```
In [15]: s_time = time.time()
```


```
W_batch, _, J_his_batch, W_his_batch = batch_gradient_descent_func(x_batch, y_batch, w,
print ('final loss : {} with {:.3f}s'.format(J_his_batch[-1], time.time()-s_time))
```

```
#print(J_his_batch[-3:])
```

final loss : 0.01871597517093907 with 0.053s

Since we only use 20 data for practice, the loss curve drops dramatically in few iterations

```
In [16]: plotCostOpt(J_his_batch)
```



To do:

- Implement `stochastic_gradient_descent_func` in `HW_YourAnswer.py` file.
- Due to randomness, the returned loss may be slightly different from the loss described in the script.

```
In [17]: # Initialize weights and biases
W = np.zeros((3072, 10))
b = np.zeros((10))
```

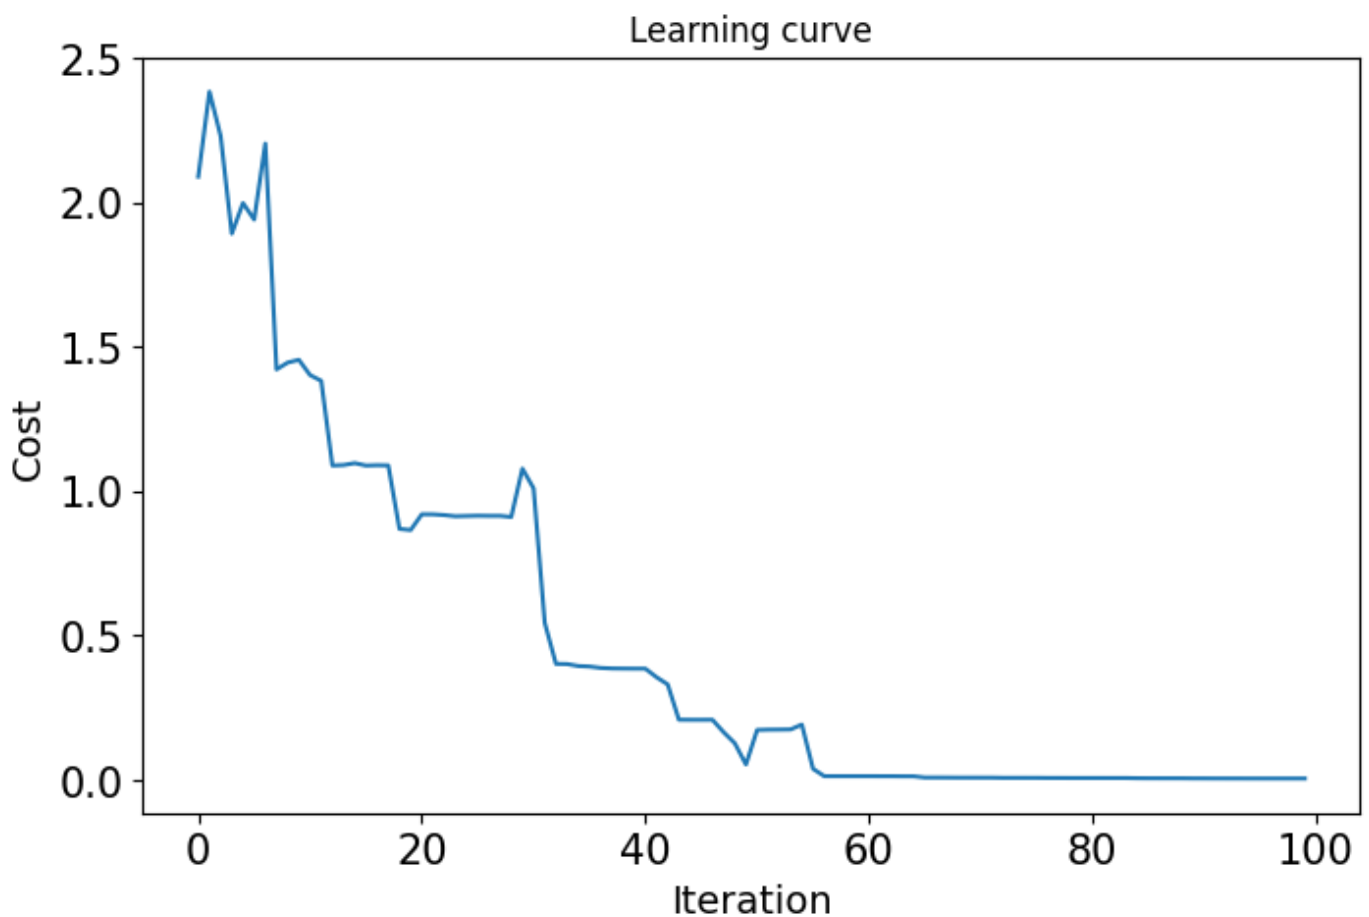
- Final loss should be 0.005614035128227841 (It doesn't have to be exactly the same)

```
In [22]: random_seed = 7
mini_batch = 2

W_SGD, _, J_his_SGD, W_his_SGD = stochastic_gradient_descent_func(x_batch, y_batch, W, b)
print ('final loss : {} with {:.3f}s'.format(J_his_SGD[-1], time.time()-s_time))

final loss : 0.005614035128227819 with 68.736s
```

```
In [23]: plotCostOpt(J_his_SGD)
```



Visualize the learned classifier

Let's see the performance of our linear classifier trained with all the training data. Also, by rescaling the weight parameters, we can visualize the learned weights for each class

```
In [24]: # Now use all the data to train our linear classifier

# 1. Training dataset
train_X = X_tr
train_Y = Y_tr_onehot
```

```

# 2. Test dataset
test_X = X_te
test_Y = Y_te

print ('Train data shape : %s, Train labels shape : %s' % (train_X.shape, train_Y.shape))
print ('Test data shape : %s, Test labels shape : %s' % (test_X.shape, test_Y.shape))

# 3. Training details
W = np.zeros((3072, 10))
b = np.zeros((10))
alpha = 0.01
iterations = 3000
mini_batch = 10

# 4. Train
W_SGD, b_SGD, J_his_SGD, W_his_SGD = stochastic_gradient_descent_func(train_X, train_Y,

Train data shape : (10000, 3072), Train labels shape : (10000, 10)
Test data shape : (1000, 3072), Test labels shape : (1000, 10)

```

- Using `linear_pred` which should be already implemented, let's evaluate on test data
- Test acc should be about 0.383

```

In [25]: pred = np.argmax(linear_pred(test_X, W_SGD, b_SGD), axis=1)
print('Test acc : ', np.sum(pred==test_Y)/len(pred))

Test acc : 0.383

```

Let's visualize the learned weight of each classes

Since our model is not that accurate, these may not look nice

```

In [26]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                        'dog', 'frog', 'horse', 'ship', 'truck']

w = W_SGD[:, :].reshape(32, 32, 3, 10)
w_min, w_max = np.min(w), np.max(w)

images_index = np.int32(np.round(np.random.rand(10,)*10000,0))


w_min, w_max = np.min(w), np.max(w)

fig, axes = plt.subplots(2, 5, figsize=(18, 6),
                        subplot_kw={'xticks': [], 'yticks': []})

fig.subplots_adjust(hspace=0.3, wspace=0.05)
# print(w_min, w_max, visualize_theta)

for idx, ax in enumerate(axes.flat):
    wimg = 255.0 * (w[:, :, :, idx].squeeze() - w_min) / (w_max - w_min)
    img = (wimg.reshape(32, 32, 3))
    ax.imshow(img.astype('uint8'))
    ax.set_title(class_names[idx])

```

Exercise 2: Neural Network Modules (Numpy)

In this exercise, we will learn to implement a Twolayer Neural Network using numpy modules. The goal here is to learn how to implement forward/backward propagation of each operations used in Neural Network and how to modularize each operations.

You will be asked to complete several functions/classes used below in `HW_YourAnswer.py`.

- `OutputLayer`
- `ReLU`
- `Sigmoid`
- `Affine`
- `TwoLayerNet`

1. Output Layer


Now let's consider output layer that forwards cross entropy loss with given score and target in class.

To do:

- Implement `OutputLayer` in `HW_YourAnswer.py` file

In [27]: `Image('fig/Output_Layer.png')`

Out[27]:



In [28]: `outputlayer = OutputLayer()`

- Forward output should be 13.095060867144417
- Backward output should be [0.90887517 -0.99999795 0.09112277]

```
In [29]: temp_x1 = np.array([[3, -10, 0.7]])
temp_t1 = np.array([[0,1,0]])
output_forward1 = outputlayer.forward(temp_x1, temp_t1)
output_backward1 = outputlayer.backward()
print('Forward propagation of output layer :', output_forward1)
print('Backward propagation of output layer :', output_backward1)
```

Forward propagation of output layer : 13.095060867144417
Backward propagation of output layer : [[0.90887517 -0.99999795 0.09112277]]

- Forward output should be 7.075548386844261
- Backward output should be [3.02958391e-01, -3.33332649e-01, 3.03742579e-02],
[-3.32509126e-01, 3.32509088e-01, 3.74189683e-08],
[7.26173786e-04, 2.92959414e-01, -2.93685588e-01]]

```
In [30]: temp_x2 = np.array([[3, -10, 0.7],[9,15,-1],[-5,1,-1]])
temp_t2 = np.array([[0,1,0],[1,0,0],[0,0,1]])
output_forward2 = outputlayer.forward(temp_x2, temp_t2)
output_backward2 = outputlayer.backward()
print('Forward propagation of output layer :', output_forward2)
print('\nBackward propagation of output layer :', output_backward2)
```

Forward propagation of output layer : 7.075548386844261

Backward propagation of output layer : [[3.02958391e-01 -3.33332649e-01 3.03742579e-02]
[-3.32509126e-01 3.32509088e-01 3.74189683e-08]
[7.26173786e-04 2.92959414e-01 -2.93685588e-01]]

2. ReLU

To do:

- Implement ReLU in HW_YourAnswer.py file

$$ReLU(x) = \max(0, x)$$

```
In [31]: relu = ReLU()
```

- Forward propagation should be [3. 0. 0.7]
- Backward propagation should be [-10 0 0]

```
In [32]: temp_x1 = np.array([[3, -10, 0.7]])
temp_x2 = np.array([[ -10,1,0]])
relu_forward1 = relu.forward(temp_x1)
relu_backward1 = relu.backward(temp_x2)
print('Forward propagation of ReLU :', relu_forward1)
print('Backward propagation of ReLU :', relu_backward1)
```

Forward propagation of ReLU : [[3. -0. 0.7]]
Backward propagation of ReLU : [[-10 0 0]]

- Forward propagation should be

$$\begin{bmatrix} 3. & 0. & 0.7 \\ 9. & 15. & 0. \\ 0. & 1. & 0. \end{bmatrix}$$
- Backward propagation should be

$$\begin{bmatrix} 3 & 0 & -10 \\ 5 & -4 & 0 \\ 0 & -5 & 0 \end{bmatrix}$$

```
In [33]: temp_x3 = np.array([[3, -10, 0.7],[9,15,-1],[-5,1,-1]])
temp_x4 = np.array([[3,5,-10],[5,-4,2],[-3,-5,3]])
relu_forward2 = relu.forward(temp_x3)
relu_backward2 = relu.backward(temp_x4)
print('Forward propagation of ReLU :', relu_forward2)
print('\nBackward propagation of ReLU :', relu_backward2)
```

```
Forward propagation of ReLU : [[ 3. -0.  0.7]
 [ 9. 15. -0. ]
 [-0.  1. -0. ]]
```

```
Backward propagation of ReLU : [[ 3  0 -10]
 [ 5 -4  0]
 [ 0 -5  0]]
```

3. Sigmoid

To do:

- Implement `Sigmoid` in `HW_YourAnswer.py` file

$$\sigma(x) = \frac{1}{1 + \exp^{-x}}$$

```
In [34]: sigmoid = Sigmoid()
```

- Forward propagation output should be [9.52574127e-01 4.53978687e-05 6.68187772e-01]
- Backward propagation output should be [0.13552998 -0.00045396 0.15519901]

```
In [35]: temp_x1 = np.array([[3, -10, 0.7]])
sigmoid_forward1 = sigmoid.forward(temp_x1)
sigmoid_backward1 = sigmoid.backward(temp_x1)
print('Forward propagation of sigmoid :', sigmoid_forward1)
print('Backward propagation of sigmoid :', sigmoid_backward1)
```

```
Forward propagation of sigmoid : [[9.52574127e-01 4.53978687e-05 6.68187772e-01]]
Backward propagation of sigmoid : [[ 0.13552998 -0.00045396  0.15519901]]
```

- Forward propagation output should be
[9.52574127e-01 4.53978687e-05 6.68187772e-01],
[9.99876605e-01 9.99999694e-01 2.68941421e-01],
[6.69285092e-03 7.31058579e-01 2.68941421e-01]

- Backward propagation output should be
[1.35529979e-01 -4.53958077e-04 1.55199011e-01],
[1.11041415e-03 4.58853200e-06 -1.96611933e-01],
[-3.32402834e-02 1.96611933e-01 -1.96611933e-01]

```
In [36]: temp_x2 = np.array([[3, -10, 0.7],[9,15,-1],[-5,1,-1]])
sigmoid_forward2 = sigmoid.forward(temp_x2)
sigmoid_backward2 = sigmoid.backward(temp_x2)
print('Forward propagation of sigmoid :',sigmoid_forward2)
print('\nBackward propagation of sigmoid :',sigmoid_backward2)
```

```
Forward propagation of sigmoid : [[9.52574127e-01 4.53978687e-05 6.68187772e-01]
 [9.99876605e-01 9.99999694e-01 2.68941421e-01]
 [6.69285092e-03 7.31058579e-01 2.68941421e-01]]
```

```
Backward propagation of sigmoid : [[ 1.35529979e-01 -4.53958077e-04 1.55199011e-01]
 [ 1.11041415e-03 4.58853200e-06 -1.96611933e-01]
 [-3.32402834e-02 1.96611933e-01 -1.96611933e-01]]
```

4. Affine

To do:

- Implement `Affine` in `HW_YourAnswer.py` file
- **Note** : bias is added separately.

$$Affine(W, b) = XW + b$$

- Note that the matrix multiplication are implemented in XW not WX which are different from the Lecture slides. Also this time, bias is considered separately (not included in the weight)
- Forward propagation output should be
[0.51 -0.39 0.84]
[-0.07 -0.02 0.02]
- Backward propagation output should be
[-0.61 0.28]
[-0.25 -0.21]

```
In [37]: temp_W = np.array([[0.2, -0.3, 0.6], [-0.9, 0.1, -0.4]])
temp_b = np.array([[0.2, -0.3, 0.6]])
temp_x = np.array([[0.2, -0.3], [-0.9, 0.1]])
temp_t = np.array([[0.1, 0.5, -0.8], [0.4, 0.7, -0.2]])

affine = Affine(temp_W, temp_b)
affine_forward1 = affine.forward(temp_x)
affine_backward1 = affine.backward(temp_t)
print('Forward propagation of Affine :\n', affine_forward1)
print('\nBackward propagation of Affine :\n', affine_backward1)
```

```
Forward propagation of Affine :
[[ 0.51 -0.39 0.84]
 [-0.07 -0.02 0.02]]
```

Backward propagation of Affine :
[[-0.61 0.28]
[-0.25 -0.21]]

- dW of affine should be
[[-0.34, -0.53, 0.02],
[0.01, -0.08, 0.22]]
- db of affine should be
[0.5, 1.2, -1.]

```
In [38]: dw = affine.dw  
db = affine.db  
print('Gradient of the weights :\n', dw)  
print('\nGradient of the biases :', db)
```

```
Gradient of the weights :  
[[-0.34 -0.53  0.02]  
[ 0.01 -0.08  0.22]]
```

```
Gradient of the biases : [ 0.5  1.2 -1. ]
```

5. TwoLayerNN

To do:

- Implement `TwoLayerNet` in `HW_YourAnswer.py` file
- Note : Our TwoLayerNN will be using L2-regularization to prevent overfitting.
- Use `OrderedDict` to make a model (<https://pymotw.com/2/collections/ordereddict.html>)

Numerical gradient vs Backpropagation

- Running time of `grad_backprop` should be much faster than the time of `grad_numerical`

```
In [39]: network = TwoLayerNet(input_size=3072, hidden_size=10, output_size=10, regularization =  
  
x_batch = X_tr[:20]  
y_batch = Y_tr_onehot[:20]  
  
start_time = time.time()  
grad_backprop = network.gradient(x_batch, y_batch)  
print("[grad_backprop] running time(sec) : " +str(time.time() - start_time))  
  
start_time = time.time()  
grad_numerical = network.numerical_gradient(x_batch, y_batch)  
print("[grad_numerical] running time(sec) : "+str(time.time() - start_time))  
  
[grad_backprop] running time(sec) : 0.002269268035888672  
[grad_numerical] running time(sec) : 85.84250950813293
```

Note that the difference is trivial

```
In [40]: for key in grad_numerical.keys():  
diff = np.average( np.abs(grad_backprop[key] - grad_numerical[key]) )  
print(key + ":" + str(diff))
```

```
W1:7.560013093594374e-11
b1:3.5416817148838966e-11
W2:1.0421137166746752e-10
b2:5.76673855109261e-10
```

Let's train our TwoLayerNet

```
In [7]: input_size_=3072
        hidden_size_=1024
        output_size_=10
        regularization_ = 0.0001

        network = TwoLayerNet(input_size=input_size_, hidden_size=hidden_size_, output_size=output_size_)

        # NOTE: you may change iters_num while debugging your code
        # but before submission, please submit the final model trained with 2000 iterations.
        iters_num = 2000
        train_size = X_tr.shape[0]
        batch_size = 100
        learning_rate = 0.1

        train_loss_list_two = []
        train_acc_list_two = []
        test_acc_list_two = []

        iter_per_epoch = max(train_size / batch_size, 1)
```

```
In [8]: for i in range(iters_num):

        batch_mask = np.random.choice(train_size, batch_size)
        x_batch = X_tr[batch_mask]
        y_batch = Y_tr_onehot[batch_mask]

        _ = network.gradient(x_batch, y_batch)

        network.update_params(learning_rate)

        loss = network.loss(x_batch, y_batch)
        train_loss_list_two.append(loss)


        if i % iter_per_epoch == 0:
            train_acc = network.accuracy(X_tr, Y_tr_onehot)
            test_acc = network.accuracy(X_te, Y_te_onehot)
            train_acc_list_two.append(train_acc)
            test_acc_list_two.append(test_acc)

            print("Epoch : ", i / iter_per_epoch + 1, "Training acc : ", round(train_acc, 2),
```

```
Epoch : 1.0 Training acc : 0.15 Test acc : 0.17
Epoch : 2.0 Training acc : 0.35 Test acc : 0.35
Epoch : 3.0 Training acc : 0.4 Test acc : 0.37
Epoch : 4.0 Training acc : 0.44 Test acc : 0.39
Epoch : 5.0 Training acc : 0.46 Test acc : 0.42
Epoch : 6.0 Training acc : 0.49 Test acc : 0.42
Epoch : 7.0 Training acc : 0.51 Test acc : 0.44
Epoch : 8.0 Training acc : 0.53 Test acc : 0.46
Epoch : 9.0 Training acc : 0.55 Test acc : 0.45
Epoch : 10.0 Training acc : 0.58 Test acc : 0.46
Epoch : 11.0 Training acc : 0.59 Test acc : 0.47
Epoch : 12.0 Training acc : 0.61 Test acc : 0.47
Epoch : 13.0 Training acc : 0.64 Test acc : 0.47
Epoch : 14.0 Training acc : 0.65 Test acc : 0.49
Epoch : 15.0 Training acc : 0.67 Test acc : 0.49
Epoch : 16.0 Training acc : 0.69 Test acc : 0.46
Epoch : 17.0 Training acc : 0.69 Test acc : 0.48
```

Epoch : 18.0 Training acc : 0.71 Test acc : 0.46
Epoch : 19.0 Training acc : 0.74 Test acc : 0.48
Epoch : 20.0 Training acc : 0.75 Test acc : 0.47

In [9]: `model_plot(train_acc_list_two, test_acc_list_two)`



In []: